

## TCSS 564

## Name- Blaise Dmello

Schema For the Analysis:

Address(A\_ID, Apt\_house\_no, Street\_Name, City, Zipcode)

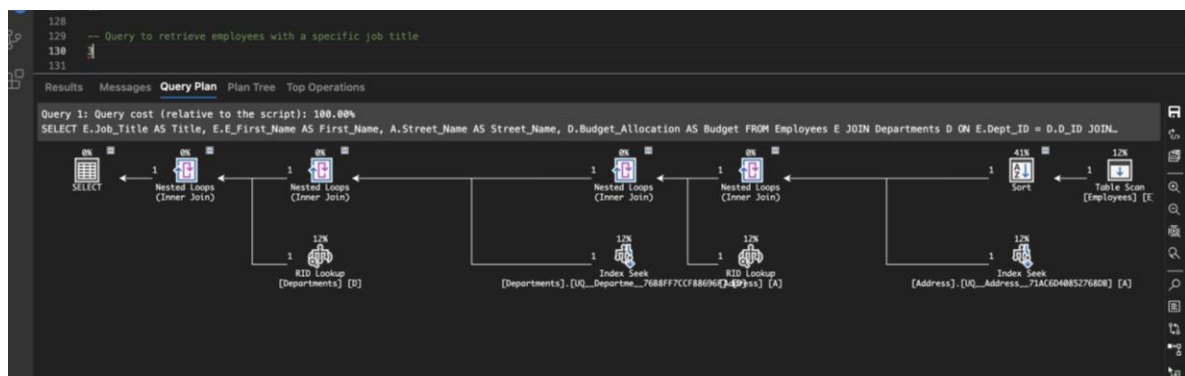
Departments(D\_ID, D\_Name, Office\_Sec, Budget\_Allocation)

Employees(EID, E\_First\_Name, E\_Last\_Name, Email, Phone, Job\_Title, Dept\_ID, Address\_ID, Salary)

### Query To Be Tested:

```
SELECT
    E.Job_Title AS Title,
    E.E_First_Name AS First_Name,
    A.Street_Name AS Street_Name,
    D.Budget_Allocation AS Budget
FROM
    Employees E
JOIN
    Departments D ON E.Dept_ID = D.D_ID
JOIN
    dbo.Address A ON E.Address_ID = A.A_ID
WHERE
    E.Job_Title = 'Software Engineer' OR
    D.D_ID < 5 OR
    A.Street_Name LIKE 'M%' OR
    D.Budget_Allocation > 1000000
ORDER BY
    E.E_First_Name ASC;
```

### 1. Initial Query Plan.



Here We can see the Query plan for the above-mentioned query on a plain table with no indexes although we haven't explicitly mentioned it, the ID fields of each of these table have a UNIQUE Constraint which acts as a form of index thus we can see that 2 of the iterators having index seeks instead of Table Scans. Which is due to the fact that scans take up  $O(n)$  time/operations whereas seeks take  $O(\log(n))$  time/operations. We are also performing RID lookups on tables Departments and Address, this is due to the lack of indexes on those tables and hence there not a B+ tree constructed being any particular order they are arranged in so seeking on indexes with the RID becomes the viable scenario in this instance.

### 2. After Creating Clustered Indexes On Each Of The Tables (i.e. Primary Keys).

```
-- Adding Primary Key Constarints to Address Table
ALTER TABLE dbo.Address
ADD CONSTRAINT PK_Address PRIMARY KEY (A_ID);
GO

-- Adding Primary Key Constraint to Departments Table
ALTER TABLE Departments
ADD CONSTRAINT PK_Departments PRIMARY KEY (D_ID);
GO

-- Adding Primary Key Constraint to Employees Table
ALTER TABLE Employees
ADD CONSTRAINT PK_Employees PRIMARY KEY (EID);
GO

-- Linking Employees Table and Address Table Via Foreign Key In Employees Table
ALTER TABLE Employees
ADD CONSTRAINT FK_Address
FOREIGN KEY (Address_ID)
REFERENCES dbo.Address(A_ID);
GO

-- Linking Employees Table and Departments Table Via Foreign Key In Employees Table
ALTER TABLE Employees
ADD CONSTRAINT FK_Department
FOREIGN KEY (Dept_ID)
REFERENCES Departments(D_ID);
GO

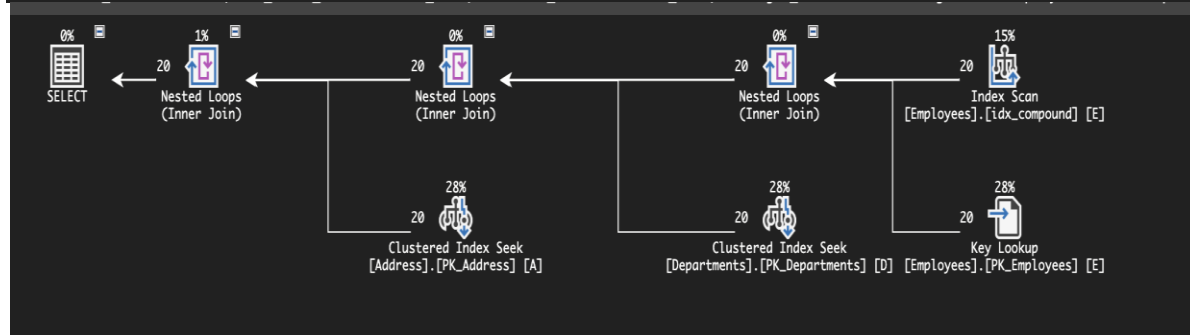
SELECT E.Job_Title AS Title, E.E_First_Name AS First_Name, A.Street_Name AS Street_Name, D.Budget_Allocation AS Budget FROM Employees E JOIN D
-- Query Execution Plan Diagram
```

Above we've altered the tables to have primary keys which in turn creates Clustered Indexes on each of the IDs. The right most iterator which was previously a Table scan changed to a Clustered index scan as the Employees table will be physically sorted in accordance with the index created on E\_ID. Furthermore, the 2 iterators at the bottom operating on Departments and Adress tables switched to a clustered index seek while also loss of the RID lookup iterators from the previous query plan. This is due to the fact that there is a clustered index physically sorting the Address and Departments table in ASC order and hence there being no need for a lookup just a Clustered Index Seek of  $O(\log(n))$  operation/time.

### 3. After Creating Non-Clustered Indexes On Employees Tables Column E\_First\_Name.

```
-- Create Unclustered Index On Employees Table
```

```
CREATE INDEX idx_compound ON Employees(E_First_Name);
GO
```



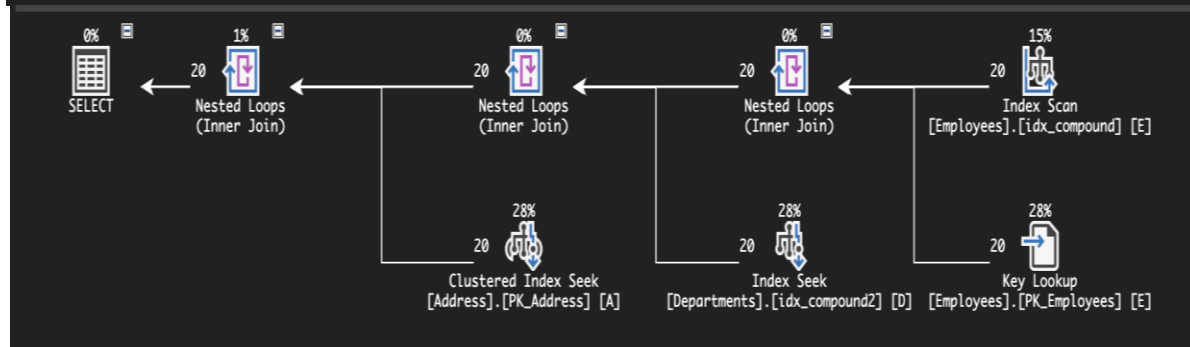
After creating a non-clustered index on Employees table, we can see that the clustered index scan from previous query plan changed to non-clustered index scan and a key lookup. This is due to the fact that scanning the entire Employees table is expensive given that it has a lot of columns. The fact that the query also doesn't require doesn't help its cause, so it makes sense that the Query optimizer chose to do a non-clustered index scan to select limited no. Of columns to reduce page/disk accesses and looking up any additional missing column required using the key lookup iterator.

#### 4. After Creating Non-Clustered Indexes On Departments Tables Columns (D\_ID, Budget\_Allocation).

```
-- Create Unclustered Index On Departments Table
```

```
CREATE INDEX idx_compound2 ON Departments(D_ID,Budget_Allocation);
```

```
GO
```



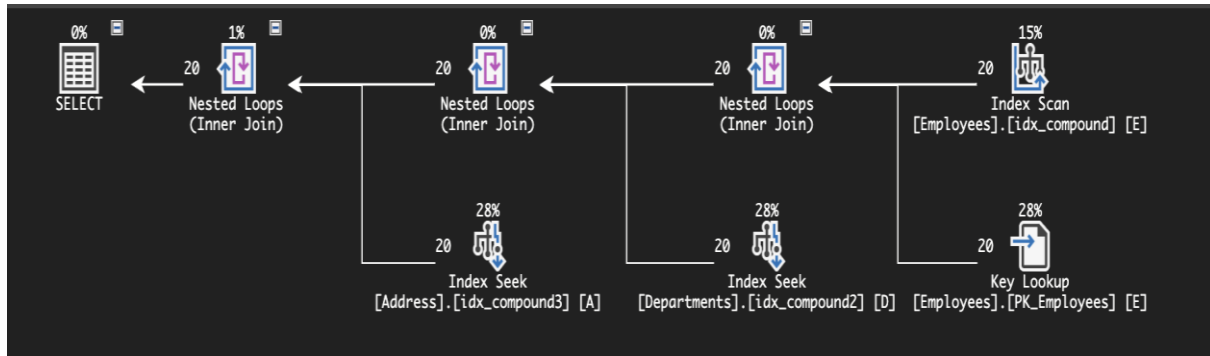
Additionally, creating non clustered index on Departments table changes the respective iterator from Clustered Index seek to just index seek. This is due to the fact that even though both are  $O(\log(n))$  operation/time the 'n' of clustered index which we will denote as N is greater than the 'n' of non-clustered index i.e.  $N > n$ . This is due to the extra unnecessary columns accessed in the clustered index thus requiring more resources.

#### 5. After Creating Non-Clustered Indexes On Address Tables Columns (A\_ID, Street\_Name).

```
-- Create Unclustered Index On Address Table
```

```
CREATE INDEX idx_compound3 ON dbo.Address(A_ID,Street_Name);
```

```
GO
```



Here we can see a similar situation as before, the iterator that previously performing clustered index seek on Address table switched to non-clustered index seek due to the size difference of  $N > n$  and thus reducing the operation and time complexity from  $O(\log(N))$  to  $O(\log(n))$  where  $N > n$  due to the size of the columns and disk accesses required.

### Estimated SubTree Cost & Elapsed Time

	Subtree Costs	Elapsed Time (ms)
<b>No Indexes</b>	<b>0.0836</b>	<b>1ms</b>
<b>Clustered Indexes on Each of the Tables</b>	<b>0.074</b>	<b>0.3ms</b>
<b>Non-Clustered Index on Employees</b>	<b>0.054</b>	<b>0.02ms</b>
<b>Non-Clustered Index on Department</b>	<b>0.077</b>	<b>1ms</b>
<b>Non-Clustered Index on Address</b>	<b>0.109</b>	<b>5ms</b>

